

AI and Modern Mercantilism

Forecasts, Framework, and Analytical Appendix

My central thesis: Modern Mercantilism will durably persist as an institutional structure. Strategic competition over compute, semiconductors, and critical minerals will prolong the AI-capex cycle, even as financial stress emerges among leveraged infrastructure providers.

Prepared by *Kiernan DiMeglio* · August 1, 2026

Forecasts

Ten binary forecasts on Modern Mercantilism, AI, and their intersection.

I. Modern Mercantilism Entrenches

1. Executive authority remains the foundation of the U.S. tariff regime

There is an **85% chance** that, for at least 60 days during Q4 2027, U.S. tariffs imposed under executive-delegated authority (Sections 201, 232, 301, or 338, or another such authority) will apply to imports representing at least 10% of 2026 U.S. goods-import value while Congress has enacted no new law directly establishing a generally applicable tariff schedule of comparable ($\geq 10\%$ of 2026 import value) scope.

- *Resolves YES if* both conditions hold for at least 60 days between October 1 and December 31, 2027. Temporary injunctions or administrative pauses lasting fewer than 32 days do not independently produce a NO.

2. The U.S. effective tariff rate remains above 5%

There is an **87% chance** that total U.S. customs duties collected will equal at least 5% of the customs value of U.S. goods imports in each calendar year from 2026 through 2028.

- *Resolves YES* only if the annual effective rate (customs duties collected $\div$ customs value of goods imports) is at least 5.00% in all three calendar years; a rate below 5.00% in any single year resolves NO.

3. At least three additional G20 members expand China-covering economic controls

There is a **72% chance** that, between August 1, 2026 and December 31, 2028, at least three G20 members lacking a qualifying measure at the July 31, 2026 baseline will enact or materially expand legally binding outbound-investment screening (semiconductors, AI, quantum, or advanced computing) or export controls (advanced-compute hardware, AI model weights, or cloud-based AI training services) in whose practical operation China is a covered or targeted destination.

- *Resolves YES if* "material expansion" occurs — adding a covered technology category, mandatory notification requirement, transaction prohibition, or new controlled product/service. Country-neutral measures count when China falls within an operative category (e.g., "country of concern," "high-risk jurisdiction"). An EU-level measure counts once for the EU; an individual EU member counts separately only for a national measure beyond the common EU requirement.

4. The WTO Appellate Body remains unable to hear appeals

There is a **92% chance** that the WTO Appellate Body will remain unable to constitute a division capable of hearing an appeal continuously through December 31, 2027.

- *Resolves YES if* the Appellate Body itself remains unable to hear appeals through the deadline. Alternative arrangements, including the Multi-Party Interim Appeal Arbitration Arrangement, do not count as restoration.

II. Compute and Critical Minerals Become Strategic Goods

5. A Chinese frontier model is trained primarily on Chinese-designed accelerators

There is a **60% chance** that, by December 31, 2027, a PRC-headquartered laboratory will publicly release an AI model that achieves at least 95% of the leading publicly released U.S. model's normalized score on a benchmark composite fixed at submission and was trained using Chinese-designed accelerators for more than 50% of its estimated total training compute.

- *Resolves YES if:* "publicly released" means downloadable weights or broad public API access; "Chinese-designed" means the core processor architecture and controlling IP are owned by a China-headquartered entity (non-Chinese fabrication does not disqualify); the hardware condition is met by a credible laboratory disclosure or substantially consistent conclusions from at least two pre-designated independent semiconductor/AI-infrastructure research organizations. The benchmark suite, normalization method, and retired-benchmark replacement rule are frozen in the analytical appendix.

6. China retains more than 80% of magnet rare-earth processing

There is a **90% chance** that China will account for more than 80% of global magnet rare-earth separation and refining output (e.g., neodymium, praseodymium, dysprosium, terbium) in calendar year 2028.

- *Resolves YES* on processing (separation and refining), not mining, using the first authoritative dataset reporting actual or best-estimate 2028 production shares; where multiple estimates exist, use the designated primary source rather than averaging.

7. No Gulf-designed AI accelerator reaches deployment at scale

There is a **92% chance** that, through December 31, 2028, no GCC-headquartered entity will operate within a GCC country a cluster of at least 10,000 AI-training accelerators whose core processor design and controlling IP are owned by a GCC-headquartered entity.

- *Resolves YES if* no qualifying cluster exists. An acquired design counts as GCC-controlled only if a GCC-headquartered entity acquires more than 50% voting control and control of the relevant processor IP before the cluster enters operation; minority sovereign-fund investments do not count; packaging, hosting, financing, or operating imported U.S. or Chinese processors does not qualify.

III. AI Capex: Durable Supercycle, Localized Busts

8. Big Five capital investment is higher in 2028 than in 2026

There is an **80% chance** that the combined cash purchases of property and equipment reported by Alphabet, Amazon, Meta, Microsoft, and Oracle during calendar year 2028 will exceed their combined cash purchases of property and equipment during calendar year 2026.

- *Resolves YES if* the 2028 combined total exceeds the 2026 combined total. Sum the cash-capex figure reported in each company's filings for quarters within each calendar year, using each company's consistently reported cash-flow-statement definition in both periods; exclude equipment obtained through non-cash finance leases unless subsequently restated into cash capex (use the latest restated figures).

9. A leveraged AI-infrastructure provider suffers a distress event

There is a **67% chance** that, before December 31, 2028, at least one company on a pre-registered list (frozen at submission) of non-hyperscaler AI-infrastructure providers will experience a defined financial-distress event during a calendar year in which combined Big Five cash purchases of property and equipment increase from the preceding calendar year.

- *Resolves YES if* a qualifying company experiences a qualifying distress event in a qualifying year. Eligible companies are neoclouds, GPU lessors, and data-center developers with material exposure to at least two of: funded debt/equipment financing; reliance on GPU collateral values; large power/construction commitments; significant customer concentration; negative free cash flow; dependence on continued external financing. Qualifying distress events: bankruptcy/insolvency filing; missed debt payment not cured within 30 days; distressed debt exchange; creditor-led restructuring; transfer of material equity/assets to creditors; acquisition following default or at an enterprise value below outstanding funded debt. Ordinary refinancing, voluntary asset sales, and premium acquisitions do not count.

10. At least one Big Five company enters a capex-digestion year

There is a **45% chance** that at least one of Alphabet, Amazon, Meta, Microsoft, or Oracle will report lower cash purchases of property and equipment in a fiscal year ending in 2027 or 2028 than in its immediately preceding fiscal year.

- *Resolves YES* on a single qualifying company-year. Compare each company's annual cash purchases of property and equipment using the same accounting definition in consecutive fiscal years; a reduction caused by a restatement counts only if the restated figures still show a genuine year-over-year decline.

PART TWO

Framework & Holistic Synthesis

The Strategic Stack: How AI and Modern Mercantilism Reinforce Each Other

AI and Modern Mercantilism feed into one another in a continuous loop. AI takes a set of inputs, things like critical minerals, fabrication capacity, accelerators, electricity, data, and frontier models, and turns them into strategic assets, meaning assets that hand a country real advantage when it controls them. That reclassification is what gives governments a reason to step in. They tax trade, subsidize domestic capacity, restrict access for foreign competitors, and accept economic inefficiency in return for resilience. The intervention then forces rivals to build up the same contested layers themselves, which raises overall demand for compute and minerals, which raises the stakes again. The loop drives itself, and that is why this is a system rather than another round of tariffs.

The result is not deglobalization, and it is not every country going fully on its own. It is **managed fragmentation**. Trade and investment still happen, but inside politically drawn boundaries and under government supervision, and access becomes conditional on alignment. The ten forecasts in Part 1 follow this system end to end, from how durable executive trade authority is, through the strategic chokepoints, to the financial stress that shows up at the leveraged edge of the AI buildout.

The AI value layers... layer by layer

All layers of the AI value chain have distinct chokepoints and policy instruments.

Layer	Strategic chokepoint (2024–26)	Mercantilist instrument	Anchored in
Critical minerals	China: ~91% of rare-earth refining, ~94% of permanent magnets (IEA, 2024)	Export licensing; extraterritorial controls	F6
Fabrication & equipment	US/allies (design, SME); Taiwan (leading-edge fab)	Export controls; ally coordination	F5, F7
Accelerators / compute	US: Nvidia hardware + CUDA software moat	Chip export controls; revenue-share licensing	F5, F7
Energy / grid	National utilities; nuclear and gas buildout	Permitting; sovereign power-and-compute deals	(systemic)
Models & applications	US and Chinese frontier labs	Data localization; procurement preference; weight controls	F5

The two most valuable positions in this system are not the obvious ones. Leverage does not come from the mine or the finished model, it comes from the difficult middle of the stack, meaning refining, fabrication, and the layer that ties compute and energy together, where capacity is slow and expensive to build back up. That middle is where the loop does its real work.

1. The system reinforces itself because the tools can be swapped but the goals cannot

The durability of Modern Mercantilism does not depend on any single tariff surviving a court challenge or an election. All it depends on is that governments are able to change their tools without changing what they are trying to accomplish. A tariff that gets blocked under one authority can be rebuilt through a different delegated authority, or through an export control, a procurement rule, a subsidy, or a domestic-content requirement.

The United States showed exactly this in 2026. After its emergency tariff authority was struck down, the policy did not snap back to where it started, it simply moved somewhere else, into Section 301 actions that reached around sixty economies and into a long-dormant Section 338 authority, and Congress never had to pass a tariff schedule at all. This is why the anchor forecast of the portfolio, F1, is written about

the equilibrium itself, meaning whether executive-delegated authority stays the foundation of the whole regime, rather than about any one measure, and it is why F2 tracks a structural floor underneath the realized tariff rate instead of a headline ceiling.

Two conditions let the system harden this way. The first is that outside enforcement is weak. The WTO Appellate Body still cannot hear appeals (F4), so the multilateral system is left setting baseline rules that it can no longer force a great power to follow. The share of world merchandise trade moving on most-favored-nation terms has already dropped from roughly 80% in 2024 to about 72% by early 2026 (WTO data). The second is that intervention creates its own defenders. OECD industrial subsidies across fifteen major sectors hit \$108 billion in 2024, which is about 1.3% of covered firms' revenue and the second-highest share ever recorded, and unlike the 2009 peak there was no recession to explain it (OECD, 2024). Once firms build plants and qualify for those subsidies, they turn into constituencies for the very rules that funded them. That same logic is now spreading to allied capitals as well (F3).

2. Controlled interdependence rather than autarky, and energy as the newest contested layer

Because the middle of the stack cannot be rebuilt quickly, mercantilism ends up producing partial duplication instead of a clean break. Governments lock down the layers they are able to control and keep trading for the ones they cannot reproduce yet.

Rare-earth processing is the clearest example of this. China holds only around 60% of the mining of magnet minerals, but it holds about 91% of the refining and about 94% of the magnet production (IEA, 2024), which means Western reshoring can keep expanding for years without actually taking away the leverage, because opening new mines does not build the refining and alloying ecosystems that sit downstream of them. F6 captures this, with processing dominance lasting through 2028 even as the number of project announcements keeps climbing. The compute layer mirrors the same structure from the opposite side, where F5 tests whether China's substitution effort becomes technically credible, meaning a frontier model trained mostly on domestic accelerators, and F7 tests whether Gulf sovereign-AI capital actually buys real processor independence or just pays for an imported stack.

The one layer this framework pushes above the forecast list is energy. Electricity has turned into the physical constraint that binds the whole buildout, with global data-center consumption projected to climb from around 485 TWh in 2025 toward around 950 TWh by 2030 and AI-specific load roughly tripling over that stretch (IEA), and it is now being treated as a strategic input in its own right. The grid is taking on national-security status, nuclear restarts and long-term power-purchase agreements are being written as industrial policy, and sovereign deals increasingly bundle power and compute together in the same package. Energy is the place where the mineral, chip, and capital layers physically come together, and it is where the next round of mercantilist tools is most likely to land.

3. The capex paradox, a higher floor but a more wasteful composition

Once compute counts as strategic, capital expenditure stops being only a question of return on investment and picks up a second kind of value, which is insurance against being dependent on or shut out by someone else. That is the reason the buildout keeps going even though the revenue from applications is still uncertain. The five largest technology firms spent over \$400 billion in 2025, and that number is set to rise by roughly 75% in 2026 (IEA), because for those firms underbuilding puts model leadership, cloud demand, and access to strategically important markets at risk, while sovereign investors will fund capacity below any hurdle rate a private company would accept.

This is what clears up the apparent contradiction between a cycle that is durable and strain that is clearly visible. Mercantilism does two things at the same time. It raises the floor under total investment, since strategic demand is not very sensitive to price and is partly backed by governments, and it makes the misallocation above that floor worse, since subsidies and strategic overbuild pay for capacity that ordinary market discipline never would. So the total becomes more durable even while the composition of it becomes more wasteful. Failure does not go away, it just gets pushed into the parts of the system that mercantilism does not protect.

The three capex forecasts are built around exactly this split. F8 holds that total Big Five investment is still higher in 2028 than it is in 2026, which is the floor. F10 assigns a real probability that at least one individual platform still enters a digestion year, which is company-level strain showing up even among the strong. And F9, which is the signature call of the portfolio, predicts that outright distress shows up first among the leveraged neoclouds, the GPU lessors, and the independent developers, and that it can happen in the very same year that aggregate platform capex is going up. The physical bottlenecks make

the divergence even sharper, since the IEA estimates that grid constraints on their own could delay about 20% of planned data-center capacity, and a delay is something a hyperscaler can live through when it is funding construction out of operating cash flow, while it is fatal for a developer that is servicing debt against GPU collateral that is losing value.

Conclusion

The next decade is going to be defined less by the end of globalization than by the way it gets politically reorganized. Tariffs and controls are going to be hard to unwind, because they are held in place by three forces at once, which are executive flexibility, strategic competition, and the investment constituencies that the intervention itself creates. The mineral and compute chokepoints will drive partial duplication without ever fully cutting the dependence. AI capital expenditure will stay high because both firms and states value the capacity, for economic reasons and for strategic ones.

The system that comes out of all this will be more redundant, more capital-intensive, and less efficient than the one it is replacing, but for those exact same reasons it will also be more durable than an ordinary bubble. When the correction finally comes, it will not be the abandonment of AI infrastructure. It will be a wave of local failures among the leveraged firms built up around the edges of it, happening inside a supercycle that keeps compounding above them.

Analytical Appendix

Generative Model of the Portfolio

The ten forecasts in Part 1 are not ten separate beliefs. They are ten *readings* of a small set of latent world-states, run through a measurement layer that decides whether each reading can actually be scored in the first place. This appendix lays out that generator, derives the portfolio's probabilities, dependencies, and update rules *out of* it rather than just restating them, and, most importantly, sets out what evidence would revise the **model itself**, which is a separate thing from revising any one forecast.

A. The generator: three world-states and one measurement layer

The portfolio is generated by three latent world-states and a per-forecast resolution term:

- **M — Mercantilist policy architecture.** Executive trade authority remains fungible, industrial subsidies persist, and multilateral enforcement stays weak. Prior on "entrenched through 2028": ≈ 0.87 .
- **G — Geo-technological decoupling and chokepoint durability.** Strategic goods (minerals, fabrication, accelerators) stay contested *and* remain slow and costly to reproduce. Prior: **high**, because the binding fact is the reconstruction time of complex downstream ecosystems, which is largely exogenous to policy over a 3-year horizon.
- **K — AI-capex supercycle strength.** Aggregate infrastructure investment continues. Prior on "2028 > 2026": ≈ 0.80 .
- **R_i — Resolution quality** of forecast *i*: the probability that a true-world event is *observable and scorable* under the pre-registered criteria. Orthogonal to the world; it is a property of the question, not the outcome.

Structure among the latents. M and G share a common root, which is great-power economic competition, so they move together but are not the same thing, since policy can harden while chokepoints erode, or the other way around. The causal arrow that carries the weight is **M → K**: mercantilism raises the *floor* under capex, because strategic demand is not very sensitive to price and is partly sovereign backed. Apart from that, K is driven by commercial adoption, algorithmic efficiency, and physical constraints, and M controls none of those.

The generating equation. Every forecast decomposes into a *world term* and a *resolution term*:

$$P(F_i) \approx P(\text{indicator fires} \mid \text{latent states}) \times R_i$$

This single template explains the entire probability spread. An **anchor** is a forecast whose world term is near 1 (institutional inertia or a structural chokepoint) *and* whose R_i is clean. An **edge** is a forecast where exactly one term is deliberately loose — and the analytical content is naming *which* term.

Forecast	M	G	K	Binding term (what holds it below 1)	P
F1 executive tariff architecture	••			anchor — M near-certain	.85
F2 realized tariff floor $\geq 5\%$	••			three independent annual draws must all clear	.87
F3 allied controls diffuse	•	•		R ₃ (country-neutral drafting) + a 3-member sequence	.72
F4 WTO Appellate Body frozen	•			anchor — largely exogenous inertia	.92
F5 Chinese frontier on domestic silicon		••		R₅ (hardware-mix observability) — the binding discount	.60
F6 China >80% magnet processing		••		anchor — chokepoint durability	.90
F7 no Gulf-designed accelerator		••		anchor — Nvidia/CUDA reproduction time	.92

Forecast	M	G	K	Binding term (what holds it below 1)	P
F8 Big Five capex 2028 > 2026	↑floor		●●	K uncertainty against a high 2026 base	.80
F9 leveraged provider distress			●●(+)	conjunction: K rising and a qualifying distress event	.67
F10 a Big Five digestion year			●(-)	firm-level idiosyncrasy (only 1 of 5 needed)	.45

(●● strong loading, ● moderate; ↑floor = M lifts K's floor; + / - = sign of the K loading.)

B. Reading the marginals off the generator

Anchors (F1, F2, F4, F6, F7). The world term here is close to certain, so the amount left below 100% is really just the joint probability that several independent reversals all happen at once. F2 is the cleanest way to see this. For it to come back a NO, you would need, inside the same year, a major tariff authority getting invalidated *and* a failure to replace it with another one *and* a political appetite for fast liberalization *and* broad exemptions *and* no offsetting duties turning up somewhere else. Each of those is plausible on its own, but all of them landing together inside one calendar year is not. The 87% is a statement about that pileup, not about any single tariff.

Edges (F3, F5, F9, F10), by binding term:

- **F5 — bound by resolution, not by the world.** I would put P(a Chinese lab actually trains a frontier-class model predominantly on domestic accelerators by end-2027) at roughly **0.75**. But P(that fact is *demonstrable* to the pre-registered evidentiary standard) is lower, because Chinese labs are systematically opaque about hardware mix. The committed **0.60** $\approx$ **0.75** $\times$ **R5**. F5 is the portfolio's clearest case of a forecast priced by its measurement layer rather than its physics.
- **F9 — bound by conjunction.** $P(F9) \approx P(K \text{ rises in a given year}) \times P(\text{qualifying distress} \mid K \text{ rising})$. The forecast is deliberately *not* "an AI-infra firm fails"; it is "distress appears **while** the aggregate buildout continues," which is what isolates localized fragility from a systemic bust.
- **F10 — bound by firm idiosyncrasy.** Only one of five platforms need cut cash capex in one of two years, so this is dominated by company-specific variance rather than the K factor — which is why it sits at 45% despite a strong aggregate K.

The central causal estimate. The $M \rightarrow K$ arrow can be read directly off the Part 1 joint prior:

$$P(K \mid M) = 0.72 / 0.87 \approx \mathbf{0.83} \quad P(K \mid \neg M) = 0.08 / 0.13 \approx \mathbf{0.62}$$

Mercantilist entrenchment is worth about **+21 points** to the probability that the capex cycle lasts to 2028. That single number is the quantitative claim of the framework: strategic competition is a floor underneath investment, not a guarantee of it. (It inherits the subjectivity of the 2×2 that produced it, so it should be treated as an order-of-magnitude estimate and not a measured coefficient.)

C. How the forecasts interact: correlation as shared loadings

Dependencies are not hand-assigned; two forecasts covary **if and only if** they load on a common latent factor. The correlation structure therefore falls out of the loadings table, not out of intuition:

- **M-cluster {F1, F2, F3, F4}** — positively correlated through M.
- **G-cluster {F5, F6, F7}, plus F3** — positively correlated through G; the M–G root correlation weakly links this cluster to the M-cluster (F3 sits in both, by design).
- **K-cluster {F8, F9, F10}** — the interesting one. **F8 and F9 load positively together**, which inverts the naive bubble intuition: a *stronger* supercycle attracts more debt-financed entrants, larger construction commitments, and higher GPU-collateral valuations, enlarging the fragile stock *before* the demand profile is known. F10 loads *negatively* on K, so F8–F10 is mildly negative and F9–F10 is ambiguous (both can rise if K is high but one platform idiosyncratically trims).

The distinctive empirical content of the whole model is this **correlation structure**, and it is exactly what Section F uses to test the model, because a wrong model fails in the *pattern* of the outcomes, not just in their levels.

D. The measurement layer: pinning down R before outcomes are visible

Pre-registration is not housekeeping; it is how each R_i is fixed *ex ante* so hindsight cannot silently reprice the question. Three forecasts carry the portfolio's material resolution risk.

F5 (sets R5 — the largest discount). Benchmark composite frozen at submission:

Category	Weight
Advanced reasoning / general knowledge	25%
Mathematics and scientific reasoning	25%
Software engineering	20%
Long-context / multi-step completion	15%
Agentic / tool use	15%

One primary and one same-category replacement benchmark named per row; replacement permitted only on formal retirement, documented contamination, discontinued reporting, or saturation above a pre-set threshold — never because a result is inconvenient. Hardware evidence is scored in priority order: developer disclosure → supplier/foundry documentation → agreement of ≥ 2 pre-designated independent analysts → regulatory disclosure → technically-substantiated investigative reporting. Anonymous or unsupported claims do not resolve the hardware condition.

F9 (sets R9). The eligible universe — ~15–20 neoclouds, GPU lessors, and data-center developers — is frozen at submission; a firm qualifies only with material exposure to ≥ 2 of the Part 1 criteria (funded debt, GPU-collateral dependence, large power/construction commitments, customer concentration, negative FCF, external-financing dependence). Distress is defined financially (default, uncured missed payment, distressed exchange, creditor-led restructuring, or acquisition below funded debt), not narratively. No post-submission additions count.

F3 (sets R3). Pre-commit the two contestable calls: China counts as a target whenever it falls inside an operative category ("country of concern") even in country-neutral drafting; an EU-level measure counts once, a member-state measure only if it exceeds the common EU requirement. **F6** pre-designates its primary dataset (IEA) and single fallback (USGS) rather than averaging.

E. Update rules: move the factor, not the indicator

The governing principle: revise $P(F_i)$ only insofar as evidence shifts the **posterior on a latent factor**, with magnitude proportional to how *diagnostic* the evidence is for that factor. Evidence that moves one indicator idiosyncratically should barely move its cluster-mates, because the remaining indicators anchor the shared factor.

Evidence	Factor	Diagnostic or idiosyncratic	Move
Realized annual tariff rate prints below 6% with no replacement authority	M	diagnostic (realized, not announced)	↓ F1, F2
Two qualifying G20 members adopt China-covering controls	M, G	diagnostic	↑ F3
An allied government formally abandons an outbound-screening initiative	M	partly idiosyncratic	↓ F3 (small)
Credible ≥ 2 -analyst evidence of large-scale Chinese domestic-accelerator training	G / R5	diagnostic (relaxes the measurement bind)	↑ F5
A strong Chinese model shown to rely mainly on imported chips	R5	resolution, not world	↓ F5
China's magnet-processing share falls below 85% before 2028	G	diagnostic	↓ F6, and ↓ G broadly
A GCC entity acquires controlling processor IP	G	diagnostic	↓ F7 pre-deployment
Coordinated multi-firm cut in multi-year capex guidance ($>10\%$, ≥ 2 firms)	K	diagnostic	↓ F8, ↑ F10

Evidence	Factor	Diagnostic or idiosyncratic	Move
One platform trims a single quarter	K	idiosyncratic	~no move to F8
Power/construction continues while cloud utilization weakens	K	diagnostic for fragility	↑ F9
Marginal providers refinance without dilution or concessions	K	diagnostic	↓ F9

The asymmetry between the last two capex rows is the whole point: a single firm trimming one quarter is noise on K; a *coordinated* guidance cut is signal, because it implies digestion is systemic rather than company-specific.

F. What would revise the model

A probabilistic portfolio does not get falsified by individual misses. An 85% forecast is *built* to fail about one time in seven. What falsifies the model is **structural** failure, meaning the moments when the outcomes break the *correlation pattern* the generator predicts. There are four such failures, and each one goes after a specific factor:

1. **Decorrelation within the M-cluster (the sharpest disconfirmation).** If F1 comes back YES, meaning executive authority holds up on paper, while F2, F3, and F4 all fail, meaning the realized rate collapses, allied diffusion stalls, and enforcement comes back, then "mercantilist entrenchment" was just a label with nothing latent sitting behind it. The M factor gets rejected *even though one indicator held*, because a real common cause is supposed to move its indicators together.

2. **The M → K arrow proves illusory.** If M stays entrenched but capex rolls over anyway, meaning F8 comes back NO with M still intact, then the strategic-floor mechanism, which is the central causal claim of the appendix, is wrong, and it does not matter how M and K each behave on their own.

3. **Chokepoint durability (G) collapses fast.** China's processing share falling below 85% before 2028, or the Gulf or China reproducing a contested layer within the horizon, would show the reconstruction-time assumption was overstated and G was mispriced across F5 through F7 all at once.

4. **The fragility prediction inverts.** A buildout that keeps going with *no* marginal distress in a rising-capex year, meaning F9 comes back NO, because state guarantees and sovereign contracts end up protecting the edge, would mean that mercantilism holds up not just the aggregate but the marginal firm too, and the model explicitly says that it should not.

So the real claim of the framework is a narrow one, and it can be tested. It is **not** that government intervention raises every forecast. It is that strategic competition **raises the floor beneath investment while pushing risk out toward the leveraged edge**, which is a joint claim about *levels and correlations* at the same time. The generator above is built so that both halves can be checked, and Section F names the specific worlds in which each half turns out to be false.